

Predicting Next Pitch

Springboard Capstone - 2



Iris Martinez

Table of Contents

Table of Contents	2
Introduction	3
Problem Statement	3
Audience	3
Data Source	4
Exploratory Data Analysis	5
Initial Machine Learning Model Test Run	10
Additional Feature Engineering	11
Second Machine Learning Models Test Run	13
Conclusion	15
Future Work for Potential Model Improvements	16

Introduction

Strategy in Major League Baseball is ever involving. With hitters increasingly focused on launch angles and elevating the ball, it is unsurprising that teams are becoming more reliant on home runs to score. In fact, in 2019 the teams collectively hit an eye-popping 6,776 home runs, which is roughly a 10% increase over the previous record of 6,105 home runs hit in 2017. This trend is very likely to continue in the coming seasons.

It also is generally accepted that fastballs are particularly ripe for batters to hit for home runs given their increased velocity and lack of movement relative to other pitch types. Indeed, Fantasy Labs, a website focused on providing strategic advice to fantasy sports players noted that fastballs were the most common pitch type to be hit for home runs. And simple common sense dictates that a batter will have an even greater chance of hitting a home if he knows that the pitcher is going to throw a fastball - one only needs to watch batting practice or the home run derby to see this confirmed.

Problem Statement

The goal of this project is to build a model that attempts to accurately predict when pitchers will throw a fastball. The targeted clients for this model are professional, semi-professional, and collegiate baseball teams, including the players, managers, and other personnel. The organizations could use this data to inform their batters to anticipate fastballs under certain circumstances, which in turn should lead to the team hitting more home runs. Conversely, organizations could use this information to inform their pitchers as to when fastballs are most often thrown and coach the pitchers to throw a different pitch under those circumstances. Professional sports gamblers could also use the model to inform their wagers. These are just a handful of what is likely many uses for an accurate fastball prediction model.

Datasource

Dataset:

Pitch-level data for every pitch thrown during the 2015-2018 MLB regular seasons. Data scraped from <http://gd2.mlb.com/components/game/mlb/>. Each row represents a single pitch. I am using this dataset for training the model.

<https://www.kaggle.com/pschale/mlb-pitch-data-20152018?select=pitches.csv>

Dataset Key:

CH - Changeup

CU - Curveball

EP - Eephus*

FC - Cutter

FF - Four-seam Fastball

FO - Pitchout (also PO)*

FS - Splitter

FT - Two-seam Fastball

IN - Intentional ball

KC - Knuckle curve

KN - Knuckeball

PO - Pitchout (also FO)*

SC - Screwball*

SI - Sinker

SL - Slider

UN - Unknown*

px : x-location as pitch crosses the plate. X=0 means right down the middle

pz: location as pitch crosses the plate. Z=0 means the ground

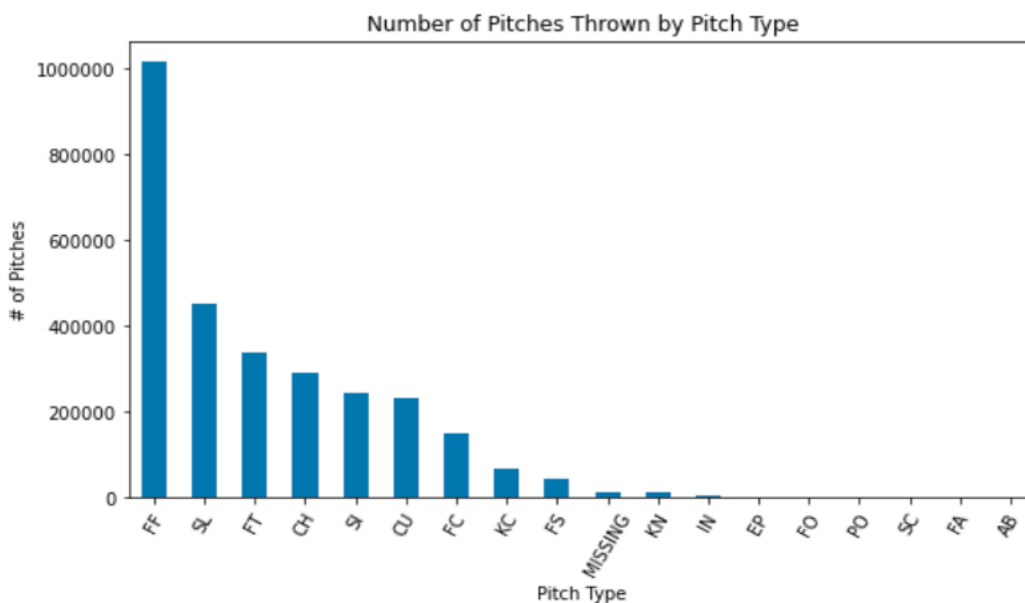
Exploratory Data Analysis

Pre-Pitch Categorical Variables


The second type of features in the dataset are pre-pitch categorical variables. These are more important for the building of my model because, unlike the post-release continuous variables discussed above; they are known prior to the pitch being thrown.

Some examples of these circumstances are runners on base, the batting count, the game score, the pitch number in the at-bat, the handedness of the pitcher and batter, and the number of outs.

In order to accurately evaluate the various categorical variables, I first needed to identify a default baseline percentage regarding how often a pitcher throws a fastball. This can be done by calculating the percentage of pitches that are labeled as a fastball in the entire dataset:




Percentage of Fastballs Thrown



As we can see, a pitcher threw a four-seam, two-seam, or cut-fastball approximately 52.41% of the time. I can use this number as a baseline when analyzing the various categorical data variables. A variable is not overly significant to the extent that it correlates to a fastball usage at or around that same 52.41% frequency - the goal is to identify the circumstances in which that frequency increases (or decreases) materially from that baseline number.


Runners on Base

The most significant influence on fastball usage is whether there is at least one runner in scoring position (on second or third base). Interestingly, however, the fastball usage spikes back up to above average when the bases are loaded, likely because a walk in those circumstances result in a run scored.



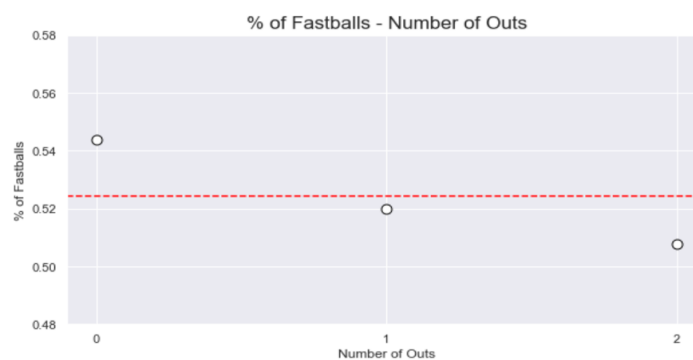
Batting Count

A pitcher's fastball usage decreases when he is ahead in the count and increases when he is behind in the count (likely because he needs to try to avoid walking the batter).



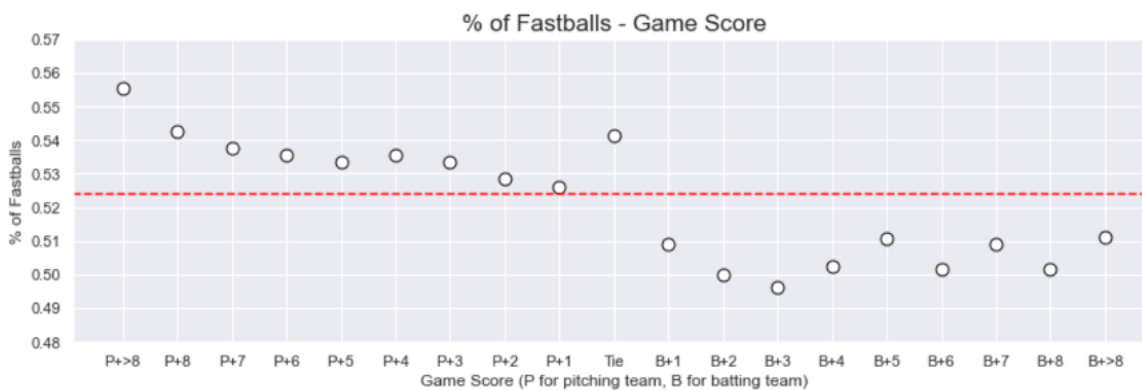
Number of Outs

Fastball usage decreases depending upon the number of outs.



Game Score

Pitchers throw more fastballs than usual when their team is winning, or the game is tied and significantly less fastballs when their team is behind.



In summary, it certainly appears as though some correlation exists between certain game circumstances and the likelihood that a pitcher will throw a fastball on the next pitch. I therefore am hopeful that a model can predict when a fastball will be thrown at a rate higher than the default 52.41 percent.

Summary

Based on the above exploratory data analysis, I have drawn the following conclusions about the dataset:

- The batting count appears to be the most significant indicators of fastballs - batters should expect to face significantly more fastballs when ahead in the count (and a full count) and significantly less when behind in the count
- The break length of a pitch is inversely correlated to fastball likelihood percentage
- A pitch with a high spin rate and spin direction at or near 180 degrees is most likely a fastball
- Pitchers are less likely to throw fastballs with runners in scoring position (unless the bases are loaded)

-
- Pitchers throw more fastballs when there are no outs in an inning
 - Pitchers throw more than an average amount of fastballs in the first, seventh, eighth, ninth, and extra innings. Pitchers throw less than an average amount of fastballs in the third, fourth, fifth, and sixth innings
 - Right-handed pitchers throw significantly more fastballs than left-handed pitchers
 - Batters should expect to face less than an average amount of fastballs when their team is winning by 4 or less runs. Pitchers will gradually throw more fastballs as the pitching team's lead increases or the batting team gets out to a 5 or more run lead

While I found variables that had some correlation with fastball likelihood percentage, no one single variable appears to be outcome-determinative of when pitchers throw fastballs. I am optimistic that this data presents circumstances to build a robust model with higher accuracy once all the variables are analyzed together.

Initial Machine Learning Model Test Run

At this point in the project, I decided to feed the dataset into four classifiers to see how they would perform. Because my dataset has over 2.8 million samples and would be a significant time investment to train the models on the full dataset, I elected to feed the classifiers approximately 5% of the data and confirmed that the sample data had the same 52/48% split between fastballs and non-fastballs.

I also used only the pre-pitch categorical features discussed in the previous section. Including the post-release continuous features would no doubt dramatically improve model performance, but would defeat the purpose of this exercise by using information not available to the batter prior to the pitch being thrown.

The initial models yielded the following accuracy on the test set from the sample data:

- Logistic Regression Classifier: 0.5392
- SGDClassifier: 0.5336
- Random Forest Classifier: 0.5240
- Gradient Boosting Classifier: 0.5484

The results were not encouraging. The classifiers performed only marginally better than the default 52.41% accuracy that you could achieve simply by guessing fastball every time. So, it appeared at least on the surface that Major League Baseball pitchers were fairly adept at mixing up their pitch selection in the various game circumstances that I fed into my models.

I had not, however, fed any information into the models regarding the previous pitch thrown by the pitcher. That information would be known by the batter and was available in my dataset, so I decided to add some additional features to the dataset fed into the models.

Additional Feature Engineering

I was confident that information about the previous pitch in a pitch sequence should improve the accuracy of my models. The major obstacle to this approach, however, was that building previous pitch features would involve having to iterate over the samples to include the data only when relevant (i.e., only to pitches that are not the first pitch of a particular at-bat), which was very time-consuming and computationally expensive given that there are over 2.8 million pitches in my dataset. I ultimately decided to add the following features:

-
- Inning Pitch Count - this feature provides how many pitches a pitcher has thrown in a particular inning. I could only tally this efficiently on an inning-by-inning basis (rather than a pitch count for the entire game) given how the data was aggregated in my dataset.
 - Previous Pitch Type - this feature provides the pitch type of the immediate previous pitch thrown to the batter using the pitch type labels (Once this data was compiled, I converted it to categorical features using dummy variables)
 - Vertical/Horizontal Previous Pitch Location - these two features provide the vertical and horizontal pitch location of the immediate previous pitch to a batter in a particular at-bat. These were built off of the 'px' and 'pz' continuous variables discussed above.

I ended up settling on using only the previous pitch type and location features due to how computationally expensive it was to build these features over a 2.8 million sample dataset. Had I not been under such a time constraint, I would have built a previous pitch feature for each of the continuous variables discussed in the exploratory data analysis section of this report. I ultimately chose to simply use the previous pitch type label knowing that the particular label somewhat aggregates that data into one feature.

Second Machine Learning Model Run

Having added the previous pitch features to the dataset, I fed the updated 5% sample data into some models to see if I had improved performance. Here are the resulting test set accuracy numbers for each model:

-
- Logistic Regression Classifier: 0.601517
 - SGDClassifier: 0.591917
 - Random Forest Classifier: 0.569545
 - Gradient Boosting Classifier: 0.601848

As we can see, adding features regarding the previous pitch in a pitch sequence significantly improved models.

Third Machine Learning Model Run

GridSearchCV is used to optimize our classifier and iterate through different parameters to find the best model. I used this on all our models to see if this would improve the models' accuracy.

- Logistic Regression Classifier: 0.656
- SGDClassifier: 0.648
- Random Forest Classifier: 0.62
- Gradient Boosting Classifier: 0.604

Hyper-parameterizing our models using GridSearchCV significantly improved our models' accuracy.

The Logistic Regression Classifier model was able to predict 65.6% of fastballs, that is over 10% better than blindly guessing that a pitcher will throw a fastball 52.41% of the time.

Random Forest Classifier Feature Importance

The Random Forest classifier only produces positive “feature importance” values to each model feature on a normalized (0-1) scale.

These values will indicate to us which features are most important to the model in making its key decisions within the decision trees, with a higher value indicating that the feature is more influential. Here are the 15 most (on the top) and least (on the bottom) features in the model:

	coefficient	rfc_feature_importance
pz_prev	-0.062737	0.157753
px_prev	-0.048638	0.156829
pitch_num	-0.032886	0.070032
stand_num	0.018800	0.029257
top_num	-0.068869	0.027969
p_throws_num	0.058830	0.025329
outs_1.0	-0.047877	0.024299
prev_fastball	0.614518	0.023597
outs_2.0	0.023413	0.022506
on_1b	0.109228	0.022212
pitcher_lead_0.0	-0.061119	0.020516
b_count_1.0	-0.015333	0.019855
on_2b	-0.070866	0.019732
s_count_1.0	-0.047185	0.017912
pitch_num_2.0	0.070412	0.016219

	coefficient	rfc_feature_importance
inning_10	-1.055783	0.005890
pitcher_lead_4.0	-0.816536	0.006904
pitcher_lead_-9.0	-0.500451	0.000821
pitcher_lead_-6.0	-0.476489	0.002827
outs_1.0	-0.386676	0.018196
pitcher_lead_5.0	-0.375364	0.005536
inning_5	-0.330019	0.013156
pitch_num_10.0	-0.314202	0.000343
s_count_1.0	-0.313586	0.017046
pitcher_lead_-5.0	-0.294217	0.002742
pitch_num_8.0	-0.281542	0.001725
s_count_2.0	-0.265491	0.014646
pitcher_lead_2.0	-0.248157	0.011091
pitcher_lead_-1.0	-0.241768	0.011143
b_count_1.0	-0.241483	0.015446

Conclusion

This project demonstrates that it is possible to analyze pre-pitch circumstances to predict whether a pitcher is going to throw a fastball as the next pitch at a higher accuracy rate than the default frequency percentage of 52.41%. It is important to note that the models struggled to outperform the default rate without information regarding the previous pitch despite observing correlations between fastball frequency and certain of these circumstances (runners on base, batting count, pitch sequence, game score, etc.).

It is interesting to note that, while correlations between certain game circumstances and fastball frequency do exist, the same overall game circumstances did not significantly improve the algorithms' ability to predict fastballs when considering all of that data in the aggregate.

Once we hyper-parameterized our features and included information about the pitcher's inning pitch count, the immediate previous pitch type thrown, and the location of the immediate previous pitch, the models' respective accuracy improved significantly by approximately 5-10%.

Teams therefore should be able to use one of the final models to correctly identify when a pitcher will throw a fastball or non-fastball more often than competitors not using such a model.

The Logistic Regression Classifier should be used for most circumstances given its overall performance. As explained in detail above, the Logistic Regression Classifier performed substantially better in identifying non-fastballs and generated higher accuracy numbers on most of the model features. It also appears to have more potential to improve through future work.

Future Work for Potential Model Improvements

I believe that I potentially could improve the performance of the final models generated by this project with some additional work. The main hurdles that I encountered with this project was the large size of the data (over 2.8 million pitches) and the limited computational power and time available. With more time and access to more computational power, I think the models could potentially be improved through exploration of one or more of the following:

- Adding more previous pitch data as additional features. I was only able to add certain previous pitch features to the dataset given how much time/computational power it took to build them. If I had more time and computational power, I would suggest adding the other previous pitch data available for the immediate previous pitch (pitch speed, break length, spin rate, etc.). I also would suggest adding the same information for each previous pitch for an entire at-bat (rather than only the data from the immediate previous pitch). I believe this has the potential to improve the model performance given how significantly the previous pitch features improved my model in this project.
- Predicting the pitches for specific pitchers. Each pitcher has their own unique pitch arsenal that they throw with unique percentages.